

RESEARCH STATEMENT

Oshani Seneviratne (senevo@rpi.edu)

December 2021

My research focuses on the novel application of decentralized information systems such as the world wide web and blockchain technologies to advance human endeavors. From the invention of the world wide web three decades ago to the formulation of blockchain technologies only a little over a decade ago, decentralized information systems have transformed various aspects of human experience from healthcare, finance, education, entertainment, research, and social connections through creative and timely breakthroughs. Various fields at the intersection and periphery of decentralized information systems, such as artificial intelligence, knowledge representation, knowledge graphs, security, privacy, and human-computer interaction, are essential components of my research.

Past Research

In my past research, I have experienced the academic, industrial, and standards-body points of view on the conceptualization, implementation, and application of decentralized information systems research within computer science.

During my graduate studies at MIT in computer science, I conducted research on policy-aware, accountable systems for the web [1] at the Decentralized Information Group under the supervision of Sir Tim Berners-Lee, the inventor of the world wide web. I developed a novel web protocol called HyperText Transfer Protocol with Accountability (HTTPA) [2], which demonstrated its effectiveness in several domains, including electronic health records transfer and creative content reuse in web-based systems [3]. HTTPA utilizes a distributed hash table-based implementation for web-based peer nodes to discover protected content expressed using a policy language grounded in semantic web technologies [4]. HTTPA allows protocol participants to declare usage restrictions and allowances of the content shared [5], and logs data access, share, reuse transactions to provide audit trails to interested rightful owners of the data items in a provenance-centric way [6]. A key innovation of this research was enabling information infrastructures to offer privacy through transparency of the data use transactions [7]. Moreover, I investigated effective tools and applications for keeping users honest through policy-aware content reuse [8]. I also co-formulated the design and development of a decentralized online social network [9], as well as an explanation interface [10] for the Tabulator¹. Some of the work leading up to my dissertation received the Yahoo! Key Scientific Challenge Award².

After graduating from MIT, I spent three years at Oracle Inc., specializing in enterprise provenance systems and governance, risk, and compliance suite of tools implemented in standards-compliant Web Ontology Language (OWL)³ rules. During this period, I had strong connections to standards bodies such as the W3C and acted as a liaison for connecting some of the standards-body work and academic research into Oracle's product offerings. To that end, I contributed to Oracle's reference implementations of the Prov Data Model⁴ and the Linked Data Platform⁵ W3C recommendations. I was also a co-inventor on a patent on visualizing provenance in enterprise information flow [11]. Furthermore, I led the research and development work on employee wellness tools incorporated into one of the core business products on human capital management.

Since joining Rensselaer Polytechnic Institute (RPI) as the Director of Health Data Research, I coordinated the research operations of the IBM Research AI-Horizon Network funded Health Empowerment through Analytics Learning and Semantics (HEALS) project⁶. I have been working very closely with the faculty principal investigators of the project on several research thrusts over the past five years. Several noteworthy research directions with respect to decentralized systems research that I led and coordinated within HEALS include: the explanation ontology for AI models in healthcare settings⁷ [12, 13], the FoodKG (a knowledge graph consisting of over 1 million recipes) [14, 15], ontological models for breast cancer diagnosis and treatment [16, 17], techniques for enabling trust in clinical decision system support recommendations through semantics [18], and personalizing health recommendations with semantics and machine learning [19, 20, 21].

Since provenance, security, privacy, policy, human-computer interaction, knowledge representation, and machine learning are critical components in effective decentralized information systems, these projects helped me cement several for future research.

Future Research

In my future research agenda, I plan to build upon my past research threads that have shown extreme promise to advance state of the art in computer science. I am interested in the following five research themes that complement each other and have the potential to revolutionize decentralized system research.

¹Tabulator is a read-write linked data browser, conceptualized by Sir. Tim Berners-Lee. For more information, please see <https://github.com/linkedata/tabulator>.

²<https://www.csail.mit.edu/news/seneviratne-selected-yahoo-key-scientific-challenges-program>

³<https://www.w3.org/OWL>

⁴<https://www.w3.org/TR/prov-dm>

⁵<https://www.w3.org/TR/ldp>

⁶<https://idea.rpi.edu/research/projects/heals>

⁷The Explanation Ontology won the best resource paper award at the International Semantic Web Conference in 2020.

(1) Interpretable AI Infrastructures: AI research has been pursuing optimal performance often at the expense of explainability and interpretability. However, the crucial questions driven by a social reluctance to accept AI-based decisions may lead to entirely new dynamics and technologies, fostering interpretability and authenticity [22]. Blockchain technologies could be used to tackle trust issues in AI models and address the lack of interpretability by providing data and cryptographically verifiable AI model provenance [23]. In particular, the following key features that are available in decentralized blockchain technologies are very desirable for interpretable AI infrastructures: (i) Transparency and visibility of the data and AI algorithms (ii) Immutability of the input data and parameters (iii) Traceability and nonrepudiation of the output (iv) Automatic execution of logic through smart contracts. In addition, blockchain technologies can enable a marketplace for data for the AI models, thus making access to the data more democratic. Owners of datasets would be rewarded for their contributions, the authenticity of the data, and other such criteria (e.g., the quality of the results of the AI models). These aspects could be automated using smart contracts for the model execution and the reward mechanism. By tracking behaviors of AI-based systems across different data input and application scenarios, we understand and have confidence in the decisions made by those systems. In case of unfortunate and unforeseen incidents that arise due to the application of the AI models, these blockchain-based “audit trails” will be essential to determine whether humans or machines (and who precisely) are at fault. Specifically, I plan to address the challenge of developing agent-driven smart contracts on the blockchain and effective information infrastructures on the web with semantics and advances in machine learning. I believe that advances in future decentralized systems would involve constructing smart contracts that act as autonomous agents with feedback, rethinking the program execution to make allowances for data ascertained at a later point in time, and the integration of various data streams for effective decision making in the absence of a centralized coordinator.

(2) Accountable Incentivized Data Sharing: I am investigating the application of decentralized techniques such as blockchain technology in addressing the challenge of designing an effective data-sharing infrastructure, with a view towards making it accountable (i.e., there is transparency in how the data is used), and the data sharing is incentivized (i.e., to ensure the adoption of the system will scale organically). I have already developed an exemplar application to address the costly and time-consuming effort needed in bringing scientific innovations from the bench (basic research) to bedside (clinical level), which has culminated in a decentralized application for researchers to share their research artifacts [24]. Since this system will operate in a decentralized environment, a particular focus will be given to investigating novel consensus mechanisms that will enable reporting the impact of scientific data and the reproducibility of the research methods, thus providing incentives to scientists in supplying well-annotated, highly reusable data and reproducible research methods. To facilitate such decentralized scientific research infrastructures, I have begun the development of an expressive, provenance-centric language, called Science Capability-based, Intention-centric, Experiment-oriented, Networked Collaborative Expression (SCIENCE) language that will capture the “science-capability” of research datasets. The SCIENCE language uses tried and tested ontologies grounded in the Semantic Web, such as the Provenance Ontology, and provides an actionable mechanism to capture provenance and responsible data usage in the scientific process. Moreover, I have analyzed how combating misinformation, such as false information on cryptocurrency, can be tackled using affordances allowed in decentralized information systems [25]. A key innovation in this proposed work is formulating a methodology for smart contracts that encodes data use agreements and shared governance of data to create computational data use agreements supporting the automatic compliance evaluation.

(3) Personalization: I am investigating the integration of personal health device data coordinated with blockchain-based technologies, specifically smart contracts, for supplementing clinical decision support systems for giving clinicians effective data points for the personalized treatment and monitoring of patients [26]. Smart contracts, at the moment, can only be implemented with explicit logic. This can be very limiting in implementing expressive logic required in complex application domains or where there is a need for plugging in machine learning models. For example, these techniques could assist in sending clinicians alerts and notifications if an adverse event is detected from the data stream from the patient’s personal health devices. In order to overcome the limitations in smart contracts, I have co-led the development of the Read-Execute-Transact-Eval-Loop (RETEL) method for smart contracts to connect to data streams from IoT devices [27]. I have also co-developed a special-purpose smart contract family called “Swarm Contracts” that have shown to be effective when there are byzantine (random and adversarial) agents with varying capabilities in decentralized settings [28]. Swarm contracts can effectively coordinate heterogeneous agents in multi-agent environments in a mutually beneficial way, optimizing personal rewards.

(4) Decentralized Edge Computing: Decentralized information systems are not complete without innovations in mobile or edge application development. The devices at the edge provide ubiquitous access to information and provide a more secure, privacy-preserving way to gain insights from a gold mine of data. There are fundamental questions on effectively collecting, organizing, performing knowledge representation, and leveraging the data from the edge. Many of these questions have not yet seen scalable answers, and given my interests and expertise, there are many potential avenues of exploration in this space. Specific to computing infrastructures for mobile application development, I have been a core research contributor to the MIT App Inventor’s *Punya* framework⁸. I have contributed to linked data components that allow a mobile application developer to quickly build applications that consume, produce, and reason over linked data [29], a framework for IoT application development [30], and building research apps leveraging the vast amounts of data available on the linked open data cloud [31]. In addition, innovations that I have co-led, which

⁸<http://punya.mit.edu>

are at the intersection of blockchain and IoT, such as BlockIoT [26], RETEL [27], and Swarm Contracts [28] provide the basis for further innovations in this space, in particular on trustworthy, self-governing, decentralized computing platforms. Many extensions of this work are possible in several different research avenues, including federated learning and personal health data insights generation.

(5) Interoperability: Last but not least, a crucial research contribution in my future research agenda will be the decentralization of semantics to allow different stakeholders to work asynchronously and separately on parts of the semantics and share the relevant pieces of the data in knowledge graphs. Safe cross-platform interactions in complex decentralized systems are rooted in the assumption that the composite parts in the system are secure and resilient to attacks. I am interested in investigating how emerging decentralized application infrastructures can be interoperable to be secure while reducing data and value silos. There is a need for robust methodologies grounded in solid computer science theory to understand systemic risks in cross-platform decentralized applications, such as the emerging Decentralized Finance (DeFi) applications. The composition of these systems is fundamentally different and will have varying degrees of security, trust, and identity mechanisms. For example, asset transfers between a financial institution to a DeFi application may have a cross-blockchain component composed of various solutions, including decentralized exchanges, money markets, and insurance. There are risks at the points of intersection that the users may not be fully aware of since the data structures, communication mechanisms, authentication, authorization procedures are different between these systems. Furthermore, having multiple blockchains performing transactions is more complicated than the sum of its parts, as there is extra complexity, and there are risks if parts of the service become obsolete and shut down, which adds many challenges to the governance of the integrated infrastructure as a whole.

Grants

I was fortunate to receive a grant as a lead PI to conduct original research, and I am currently involved in several grant proposals as outlined below.

Funding Received: I was awarded a \$100k grant in 2018 to conduct research on the Smart Contracts Augmented with Learning and Semantics (SCALES) project through the IBM-RPI Artificial Intelligence Research Collaboration. As part of SCALES, I investigated solutions for “break glass” scenarios with smart contracts, as they are inflexible to changes after deployment. The solutions I investigated utilized effective sandbox environments for smart contract execution [32], enhancing the expressivity of the decentralized application semantics using ontological concepts [33], and an exemplar healthcare application in smart contracts [34]. Furthermore, I co-investigated computational social choice mechanisms for choosing the best option for decentralized application account recovery [35].

Planned Funding Proposals: I am currently leading an NSF Smart Connected Health proposal, in which I am proposing a decentralized research data sharing platform incentivized by blockchain technologies to reduce the time it takes to bring a scientific innovation from “bench” (basic research) to “bedside” (practical application). I have also been invited to join as a co-PI or senior personnel in several other proposals, including a mega proposal that spans over 20 institutions worldwide to advance blockchain research and education funded by the Algorand Foundation, an NSF AI institute proposal on democratic group decision making, NIH R01 proposal on personal health libraries, NSF LEAP proposal on open data platform for the data-driven transformation of the building industries, an NSF STTR proposal on “Smartchain” for developing an AI-backed blockchain computing platform, and an NSF Industry-University Research Partnerships Collaboration center at RPI called Center for Research in Advances in Financial Technologies, where I have co-proposed investigating blockchain transactions for identifying “risky business” in DeFi and a framework for reducing information asymmetry powered by blockchain and IoT for the healthcare industry.

Furthermore, I am passionate about the development of and the scientific discourse in emerging research areas. I co-founded and co-organized the AICChain workshop series⁹ at the IEEE Blockchain conference in 2019, and it has been running successfully ever since. I co-chaired the Personal Health Knowledge Graph workshop series¹⁰ at the Knowledge Graph Conference in 2020 and 2021, and the AAAI Symposium on AI for Social Good¹¹ in 2019 and 2020. I am the program co-chair of the ACM Web Science conference in 2022, and I have served in the organizing and program committees of many conferences, including the International Semantic Web Conference, Web Conference, Web Science Conference, IEEE Blockchain Conference, and IEEE International Conference on Decentralized Applications and Infrastructures. I am currently a co-editor of the Semantic Technologies for Data and Algorithmic Governance issue at the Semantic Web Journal and the Personal Health Knowledge Graphs special issue at the Data Intelligence Journal. I am on the editorial boards of journals such as Web Semantics, Medical Internet Research, Social Computing, and Biomedical and Health Informatics. I plan to continue such professional service activities that I believe are mutually beneficial to the university and the research community.

In conclusion, I am optimistic that the decentralization of information is the most sustainable information model for future computing and information systems. We are seeing evidence of the success of such systems, starting from the web and the most recent widespread adoption of blockchain-based applications. There are many challenges in these nascent areas that warrant careful scientific investigation to understand and realize these technological innovations’ full potential. I am excited to be taking up these challenges and advancing the state-of-the-art in computer science and related disciplines.

⁹<https://ai4blockchain.github.io>

¹⁰<https://phkg.github.io>

¹¹<https://ai-for-socialgood.github.io>

References

- [1] **Oshani Seneviratne**. *Accountable Systems: Enabling Appropriate Use of Information on the Web*. PhD thesis, Massachusetts Institute of Technology, 2014.
- [2] **Oshani Seneviratne** and Lalana Kagal. Addressing data reuse issues at the protocol level. In *2011 IEEE International Symposium on Policies for Distributed Systems and Networks*, pages 141–144. IEEE, 2011.
- [3] **Oshani Seneviratne**. Augmenting the Web with Accountability. In *Proceedings of the 21st International Conference on World Wide Web*, pages 185–190, 2012.
- [4] **Oshani Seneviratne** and Lalana Kagal. HTTPA: Accountable HTTP. In *IAB/W3C Internet Privacy*, volume 42, 2010.
- [5] **Oshani Seneviratne** and Lalana Kagal. Usage Restriction Management for Accountable Data Transfer on the Web. In *IEEE International Symposium on Policies for Distributed Systems and Networks (IEEE Policy 2011)*, 2011.
- [6] **Oshani Seneviratne**. Data Provenance and Accountability on the Web. pages 11–24. Springer, 2021.
- [7] **Oshani Seneviratne** and Lalana Kagal. Enabling Privacy through Transparency. In *2014 Twelfth Annual International Conference on Privacy, Security and Trust*, pages 121–128. IEEE, 2014.
- [8] **Oshani Seneviratne**, Lalana Kagal, and Tim Berners-Lee. Policy-aware Content Reuse on the Web. In *International Semantic Web Conference*, pages 553–568. Springer, 2009.
- [9] Ching-man Au Yeung, Ilaria Liccardi, Kanghao Lu, **Oshani Seneviratne**, and Tim Berners-Lee. Decentralization: The Future of Online Social Networking. In *W3C Workshop on the Future of Social Networking Position Papers*, volume 2, pages 2–7, 2009.
- [10] **Oshani Seneviratne** and Tim Berners-Lee. The Point of View Axis: Varying the Levels of Explanation Within a Generic RDF Data Browsing Environment. MIT Computer Science and Artificial Intelligence Laboratory (CSAIL) Tech Report, Citeseer, 2008.
- [11] Kenneth H Beckett, Sathyamoorthy Thelungupalayam Anandan, Reza B’far, and **Oshani Seneviratne**. Visualization of Provenance Data, March 19 2019. US Patent 10,235,781.
- [12] Shruthi Chari, **Oshani Seneviratne**, Daniel M. Gruen, Morgan A. Foreman, Amar K. Das, and Deborah L. McGuinness. Explanation Ontology: A Model of Explanations for User-Centered AI. In *ISWC (2)*, volume 12507 of *Lecture Notes in Computer Science*, pages 228–243. Springer, 2020.
- [13] Shruthi Chari, **Oshani Seneviratne**, Daniel M. Gruen, Morgan A. Foreman, Amar K. Das, and Deborah L. McGuinness. Explanation Ontology in Action: A Clinical Use-Case. In *ISWC (Demos/Industry)*, volume 2721 of *CEUR Workshop Proceedings*, pages 290–295. CEUR-WS.org, 2020.
- [14] Steven Haussmann, **Oshani Seneviratne**, Yu Chen, Yarden Ne’eman, James Codella, Ching-Hua Chen, Deborah L McGuinness, and Mohammed J Zaki. FoodKG: A Semantics-driven Knowledge Graph for Food Recommendation. In *International Semantic Web Conference*, pages 146–162. Springer, 2019.
- [15] Steven Haussmann, Yu Chen, **Oshani Seneviratne**, Nidhi Rastogi, James V. Codella, Ching-Hua Chen, Deborah L. McGuinness, and Mohammed J. Zaki. FoodKG Enabled Q&A Application. In *ISWC (Satellites)*, volume 2456 of *CEUR Workshop Proceedings*, pages 273–276. CEUR-WS.org, 2019.
- [16] **Oshani Seneviratne**, Sabbir M Rashid, Shruthi Chari, James P McCusker, Kristin P Bennett, James A Hendler, and Deborah L McGuinness. Knowledge Integration for Disease Characterization: A Breast Cancer Example. In *International Semantic Web Conference*, pages 223–238. Springer, 2018.
- [17] **Oshani Seneviratne**, Sabbir M Rashid, Shruthi Chari, Jim McCusker, Kristin P Bennett, Jim Hendler, and Deborah L McGuinness. Ontology-enabled Breast Cancer Characterization. In *International Semantic Web Conference (P&ID/Industry/BlueSky)*, 2018.
- [18] **Oshani Seneviratne**, Amar K Das, Shruthi Chari, Nkechinyere N Agu, Sabbir M Rashid, Ching-Hua Chen, James P McCusker, James A Hendler, and Deborah L McGuinness. Enabling Trust in Clinical Decision Support Recommendations through Semantics. In *SeWeBMeDa@ ISWC*, pages 55–67, 2019.
- [19] **Oshani Seneviratne**, Jonathan J. Harris, Ching-Hua Chen, and Deborah L. McGuinness. Personal Health Knowledge Graph for Clinically Relevant Diet Recommendations. *American Medical Informatics Association Conference*, 2021.
- [20] Sola Shirai, **Oshani Seneviratne**, Minor Gordon, Ching-Hua Chen, and Deborah L. McGuinness. Identifying Ingredient Substitutions Using a Knowledge Graph of Food. *Frontiers Artif. Intell.*, 3:621766, 2020.
- [21] Ishita Padhiar, **Oshani Seneviratne**, Shruthi Chari, Dan Gruen, and Deborah L. McGuinness. Semantic Modeling for Food Recommendation Explanations. In *ICDE Workshops*, pages 13–19. IEEE, 2021.
- [22] Shruthi Chari, Daniel M. Gruen, **Oshani Seneviratne**, and Deborah L. McGuinness. Foundations of Explainable Knowledge-Enabled Systems. In *Knowledge Graphs for eXplainable Artificial Intelligence*, volume 47 of *Studies on the Semantic Web*, pages 23–48. IOS Press, 2020.
- [23] Shruthi Chari, Daniel M. Gruen, **Oshani Seneviratne**, and Deborah L. McGuinness. Directions for Explainable Knowledge-Enabled Systems. In *Knowledge Graphs for eXplainable Artificial Intelligence*, volume 47 of *Studies on the Semantic Web*, pages 245–261. IOS Press, 2020.
- [24] **Oshani Seneviratne** and Deborah L. McGuinness. Incentivized Research Data Sharing, Reusing and Repurposing with Blockchain Technologies. *Hawaii International Conference on System Sciences (HICSS)*, 2022.
- [25] Daniel Kazenoff, **Oshani Seneviratne**, and Deborah L. McGuinness. Semantic Graph Analysis to Combat Cryptocurrency Misinformation on the Web. In *ASLD@ISWC*, volume 2722 of *CEUR Workshop Proceedings*, pages 168–176. CEUR-WS.org, 2020.
- [26] Manan Shukla, Jianjing Lin, and **Oshani Seneviratne**. BlockIoT: Blockchain-based Health Data Integration using IoT Devices. *American Medical Informatics Association Conference*, 2021.
- [27] Manan Shukla, Jianjing Lin, and **Oshani Seneviratne**. BlockIoT-RETEL: Blockchain and IoT Based Read-Execute-Transact-Erase-Loop Environment for Integrating Personal Health Data. *IEEE Blockchain Conference*, 2021.
- [28] Jonathan Grey, Isuru S. Godage, and **Oshani Seneviratne**. Swarm Contracts: Smart Contracts in Robotic Swarms with Varying Agent Behavior. In *Blockchain*, pages 265–272. IEEE, 2020.
- [29] Weihua Li, **Oshani Seneviratne**, Evan Patton, and Lalana Kagal. A Semantic Platform for Developing Data-Intensive Mobile Apps. In *2019 IEEE 13th International Conference on Semantic Computing (ICSC)*, pages 71–78. IEEE, 2019.
- [30] Thilanka Munasinghe, Evan W Patton, and **Oshani Seneviratne**. IoT Application Development using MIT App Inventor to Collect and Analyze Sensor Data. In *2019 IEEE International Conference on Big Data (Big Data)*, pages 6157–6159. IEEE, 2019.
- [31] **Oshani Seneviratne**, William Van Woensel, Giuseppe Loseto, Floriano Scioscia, Evan W. Patton, and Lalana Kagal. Rapid Prototyping of Mobile Apps for Clinical Research using Semantic Web Technologies. In *ISWC (Posters/Demos/Industry)*, volume 2980 of *CEUR Workshop Proceedings*. CEUR-WS.org, 2021.
- [32] Shuze Liu, Farhad Mohsin, Lirong Xia, and **Oshani Seneviratne**. Strengthening Smart Contracts to Handle Unexpected Situations. In *2019 IEEE International Conference on Decentralized Applications and Infrastructures (DAPPCON)*, pages 182–187. IEEE, 2019.
- [33] Farhad Mohsin, Xingjian Zhao, Zhuo Hong, Geeth de Mel, Lirong Xia, and **Oshani Seneviratne**. Ontology Aided Smart Contract Execution for Unexpected Situations. In *BlockSW/CKG@ ISWC*, 2019.
- [34] Mengyi Li, Lirong Xia, and **Oshani Seneviratne**. Leveraging Standards Based Ontological Concepts in Distributed Ledgers: A Healthcare Smart Contract Example. In *2019 IEEE International Conference on Decentralized Applications and Infrastructures (DAPPCON)*, pages 152–157. IEEE, 2019.
- [35] Yanlin Zhu, Lirong Xia, and **Oshani Seneviratne**. A Proposal for Account Recovery in Decentralized Applications. In *2019 IEEE International Conference on Blockchain (Blockchain)*, pages 148–155. IEEE, 2019.